Session Recap: Data and Platform Concepts

DBT, repositories, and Airflow

- DBT pipelines can carry out much of the data work discussed; the opening remark about the pipelines is unclear in the transcription.
- The session estimated that code generation could cover roughly 80–90% of a job, depending on its complexity.
- The generated code belongs in the EDO data-products repository, inside its `data_products` folder.
- Airflow blueprints can be filled out with YAML to define jobs.

Databricks, Spark, and storage

- The discussion described Databricks as running DBT code by translating it into Spark.
- On the paved path, data files in S3 are expected to use **Delta format** by default.
- Delta is a columnar format. The discussion contrasted its support for data changes and consistency guarantees with the greater difficulty of inserting or deleting rows in ordinary Parquet files.
- Delta's transaction logs and manifests were noted as part of how it provides these capabilities.

Unity Catalog

- Three catalog types were identified:
  - **Application catalog**
  - **Domain catalog**
  - **Source-aligned catalog**

Data quality and materialization

- Data-quality tests can be included in data workflows.
- Materialization can be configured in two places: in the model file itself and in project-level configuration. The transcript's wording about the second location is unclear.

Topics proposed for teaching demos

- Bi-temporal loading: explain the relevant columns and one or two loading strategies.
- Harness and CI/CD were also identified as demo topics.
- The session discussion placed these topics in the Thursday-morning lineup.